

tbb::parallel_for

```
template<typename TI, typename TF>
void tbb::parallel_for( TI begin, TI end, const TF& f );
```

- Implements a parallel for loop for values in the range [begin,end)
- Each iteration **may** execute in parallel
 - There are no guarantees
- Iterations may be executed in any order
 - Entirely legal to do the iterations in opposite order to normal
- Programmer's job: make sure all iterations are independent
- TBB run-time's job: try to execute iterations as fast as possible

HPCE / dt10 / 2013 / 7.1

Implementing tbb::parallel_for

- How could we implement parallel_for?
 - Simple helper class for describing a range
- Start with a sequential version

```
template<class TI, class TF>
void parallel_for(TI beg, TI end, const TF &f)
{
    for(TI i=beg; i<end; i++){
        f( i );
    }
}
```

HPCE / dt10 / 2013 / 7.2

Implementing tbb::parallel_for

- How could we implement parallel_for?
 - Simple helper class for describing a range
- Start with a sequential version
- What about a Cilk version?
 - Q : What does the critical path look like?

```
template<class TI, class TF>
void parallel_for(TI begin, TI end, const TF &f)
{
    for(TI i=begin; i<end; i++){
        spawn f( i );
    }
}
```

HPCE / dt10 / 2013 / 7.3

Implementing tbb::parallel_for

- How could we implement parallel_for?
 - Simple helper class for describing a range
- Start with a sequential version
- What about a Cilk version?
 - Let's try again

```
template<class TI, class TF>
void parallel_for(TI begin, TI end, const TF &f)
{
    if(begin+1==end){
        spawn f( begin );
    }else{
        spawn parallel_for(begin, (begin+end)/2, f);
        spawn parallel_for((begin+end)/2, end, f);
    }
}
```

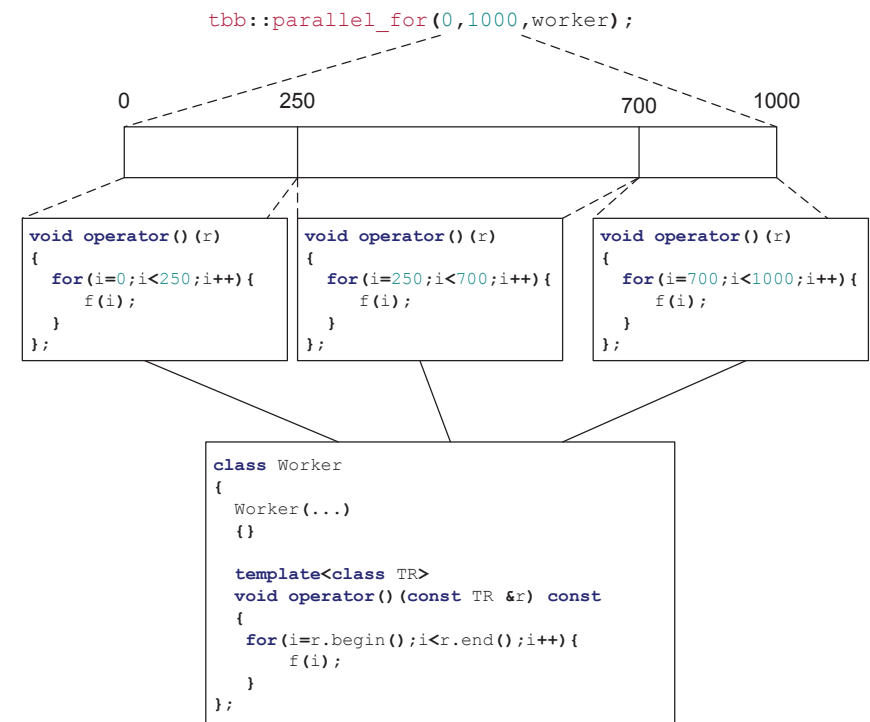
HPCE / dt10 / 2013 / 7.4

Implementing `tbb::parallel_for`

- This looks good, except what about *cost of work*!
- TBB will automatically apply **agglomeration**
 - Split range up into larger contiguous ranges
 - Hand ranges to function for processing in sequential loop
- How does it know what THRESH should be?
 - Dynamically varies based on time for last batch – auto-tune!

```
template<class TI, class TF>
void parallel_for(TI begin, TI end, const TF &f)
{
    if(end-begin < THRESH){
        for(TI i=begin; i<end; i++){
            spawn f(i);
        }
    }else{
        spawn parallel_for(begin, (begin+end)/2, f);
        spawn parallel_for((begin+end)/2, end, f);
    }
}
```

HPCE / dt10 / 2013 / 7.5



6

References vs pointers

- C++ has references as well as pointers
 - Pointers use asterisk (*), references use ampersand (&)
- Some differences between pointers and references
 - References are guaranteed to be non-null
 - A reference always points at the same object
- Operations on the references happen to the original object

```
void MyPtrFunc(my_class *x)
{
    x->wibble();
    my_class *p=x;
    *x = my_class(5);
    my_class tmp;
    x=tmp;
}
```

```
my_class x;
MyPtrFunc(&x); // pass by pointer
```

```
void MyRefFunc(my_class &x)
{
    x.wibble();
    my_class *p=&x;
    x = my_class(5);
    my_class tmp;
    &x=&tmp;
}
```

```
my_class x;
MyRefFunc(x); // pass by ref
```

HPCE / dt10 / 2013 / 7.7

Const methods

- Methods with the `const` modifier cannot change the object
 - Cannot modify the internal state of the object
 - Object is “the same” after any const method is called
- References with the `const` modifier cannot be changed
 - Can only call const methods and read object properties
 - Only methods marked const are considered const methods
 - Methods which don’t change state are “non-const” unless marked
- Const references are very useful
 - Pass by reference makes it cheap to pass object to function
 - Const-ness means object can be safely used in parallel

HPCE / dt10 / 2013 / 7.8

Scheduling through reference counts

- Each task has a reference count and a successor task
- The reference count identifies whether a task is blocked
 - If the reference count is zero then the task could be run
 - But only if it has been given to the task scheduler
 - Legal to create a task and not give it to the scheduler
 - *Note the difference: “reference count” vs “C++ reference”*
- The successor task identifies the task blocked by this task
 - Same concept as “parent” in Cilk, but slightly more general
 - When a task completes it decrements the count of its successor

HPCE / dt10 / 2013 / 7.9

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

    tbb::task * execute()
    {
        if(cond())
            return 0;
        set_ref_count(3);
        MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
        spawn(t1);
        MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
        spawn(t2);
        DoSomethingFirst();
        wait_for_all();
        DoSomethingElse();
    }
};

void CreateTasks(int start, int end)
{
    MyTask &root=new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
  
```

HPCE / dt10 / 2013 / 7.10

tbb::task	
Created	
successor	-
refcount	0
MyTask	
start	...
end	...

```

void CreateTasks(int start, int end)
{
    MyTask &root=new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
  
```

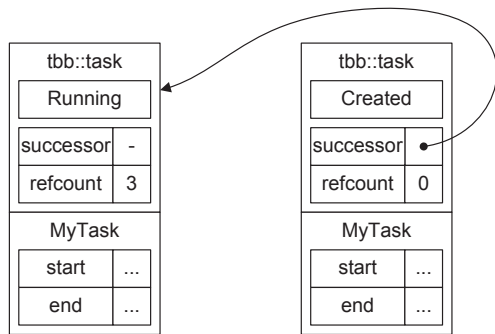
113 / 7.11

tbb::task	
Running	
successor	-
refcount	0
MyTask	
start	...
end	...

```

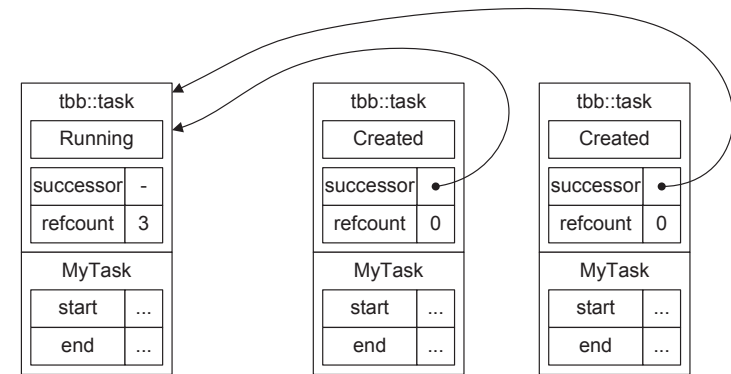
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
  
```

113 / 7.12



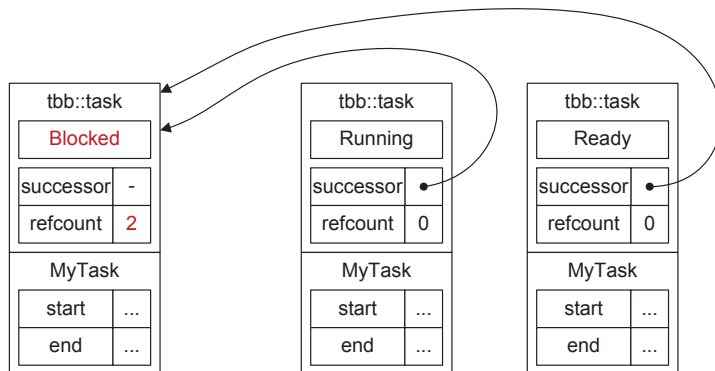
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.13



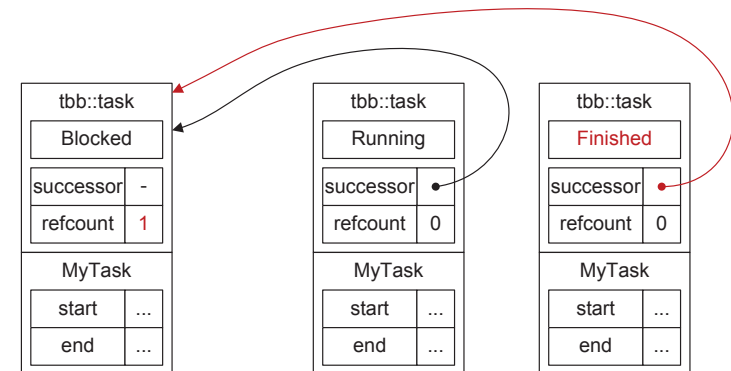
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.14




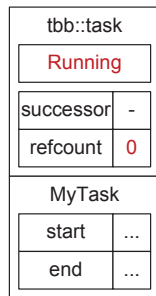
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.15




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.16




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.17



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.18



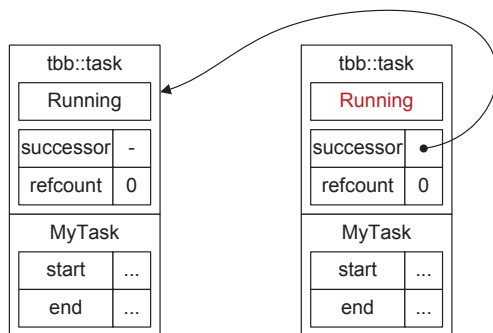
```
void CreateTasks(int start, int end)
{
    MyTask &root=new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
```

113 / 7.19

Managing reference counts

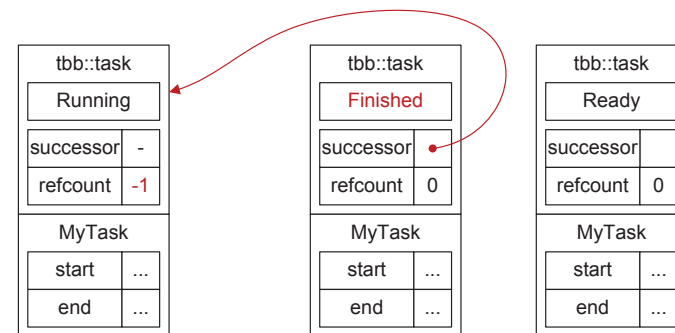
- What happens if we get the reference count wrong?
- Finishing task calls `decrement_ref_count` on successor
 - Automatically returns task to scheduler if count becomes zero

HPCE / dt10 / 2013 / 7.20



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    MyTask &t1=new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    set_ref_count(3);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.21



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    MyTask &t1=new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    set_ref_count(3);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

113 / 7.22

Some help is available

- TBB library comes in two forms: debug and release
 - release library does no error checking – all about speed
 - debug library will check reference counts at many points
- Choose library version at compilation and link stages
 - Debug: #define TBB_USE_DEBUG=1 when compiling
 - On microsoft compilers it will automatically link the correct library
 - On other compilers use “-ltbb” vs “-ltbb_debug”
 - Usually maintain different release and debug settings
 - Debug: /DTBB_USE_DEBUG=1 /MDd
 - Release: /DNDEBUG=1 /O2
 - Can setup in Visual Studio or in a makefile

Data-parallelism vs task parallelism

- Two very broad types of parallelism we’ve seen so far
- **Data-parallelism:** *do the same task lots of time in any order*
 - The code for the task stays the same for each execution
 - The input to the task changes with each execution
 - There are no dependencies between different executions
 - Often applied to elements of an array
 - Also described as “loop-parallelism”
- **Task-parallelism:** *do many different tasks with dependencies*
 - Each task has zero or more dependencies that must be met
 - Different tasks **may** have different code
 - More powerful than data-parallelism?