

# tbb::parallel\_for

```
template<typename TI, typename TF>  
void tbb::parallel_for( TI begin, TI end, const TF& f );
```

- Implements a parallel for loop for values in the range [begin,end)
- Each iteration **may** execute in parallel
  - There are no guarantees
- Iterations may be executed in any order
  - Entirely legal to do the iterations in opposite order to normal
- Programmer's job: make sure all iterations are independent
- TBB run-time's job: try to execute iterations as fast as possible

# Implementing `tbb::parallel_for`

- How could we implement `parallel_for`?
  - Simple helper class for describing a range
- Start with a sequential version

```
template<class TI, class TF>
void parallel_for(TI beg, TI end, const TF &f)
{
    for(TI i=beg; i<end; i++){
        f( i );
    }
}
```

# Implementing `tbb::parallel_for`

- How could we implement `parallel_for`?
  - Simple helper class for describing a range
- Start with a sequential version
- What about a Cilk version?
  - Q : What does the critical path look like?

```
template<class TI, class TF>
void parallel_for(TI begin, TI end, const TF &f)
{
    for(TI i=begin; i<end; i++){
        spawn f( i );
    }
}
```

# Implementing `tbb::parallel_for`

- How could we implement `parallel_for`?
  - Simple helper class for describing a range
- Start with a sequential version
- What about a Cilk version?
  - Let's try again

```
template<class TI, class TF>
void parallel_for(TI begin, TI end, const TF &f)
{
    if(begin+1==end) {
        spawn f( begin );
    }else{
        spawn parallel_for(begin, (begin+end)/2, f);
        spawn parallel_for((begin+end)/2, end, f);
    }
}
```

# Implementing `tbb::parallel_for`

- This looks good, except what about *cost of work*!
- TBB will automatically apply **agglomeration**
  - Split range up into larger contiguous ranges
  - Hand ranges to function for processing in sequential loop
- How does it know what THRESH should be?
  - Dynamically varies based on time for last batch – auto-tune!


```
template<class TI, class TF>
void parallel_for(TI begin, TI end, const TF &f)
{
    if(end-begin < THRESH){
        for(TI i=begin; i<end; i++)
            f( i );
    }else{
        spawn parallel_for(begin, (begin+end)/2, f);
        spawn parallel_for((begin+end)/2, end, f);
    }
}
```

```
tbb::parallel_for(0,1000,worker);
```

```
class Worker
{
    Worker(...)
    {}

    template<class TR>
    void operator()(const TR &r) const
    {
        for(i=r.begin();i<r.end();i++){
            f(i);
        }
    }
};
```

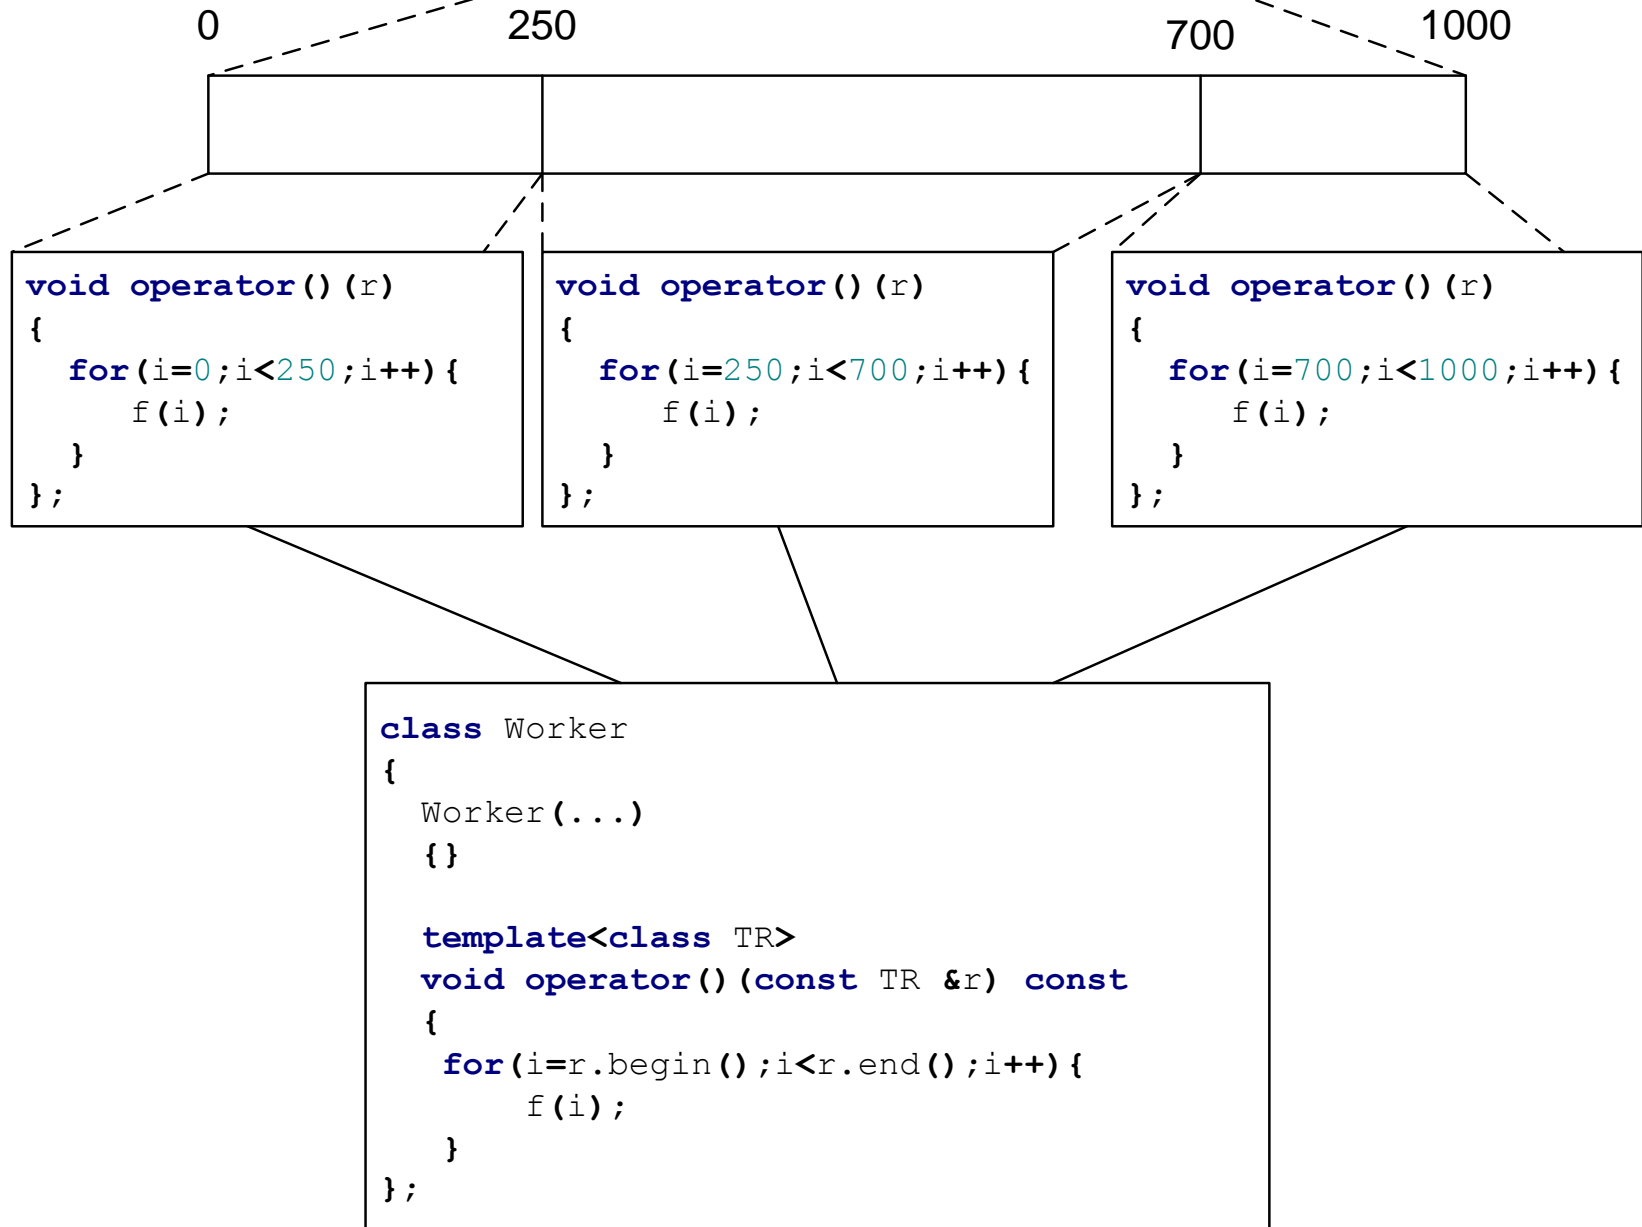
```
tbb::parallel_for(0,1000,worker);
```



```
class Worker
{
    Worker(...)
    {}

    template<class TR>
    void operator()(const TR &r) const
    {
        for(i=r.begin(); i<r.end(); i++) {
            f(i);
        }
    }
};
```

```
tbb::parallel_for(0,1000,worker);
```





# References vs pointers

- C++ has references as well as pointers
  - Pointers use asterisk (\*), references use ampersand (&)
- Some differences between pointers and references
  - References are guaranteed to be non-null
  - A reference always points at the same object
- Operations on the references happen to the original object

```
void MyPtrFunc(my_class *x)
{
    x->wibble();
    my_class *p=x;
    *x = my_class(5);
    my_class tmp;
    x=tmp;
}
```

```
my_class x;
MyPtrFunc(&x); // pass by pointer
```

```
void MyRefFunc(my_class &x)
{
    x.wibble();
    my_class *p=&x;
    x = my_class(5);
    my_class tmp;
&x=tmp;
}
```

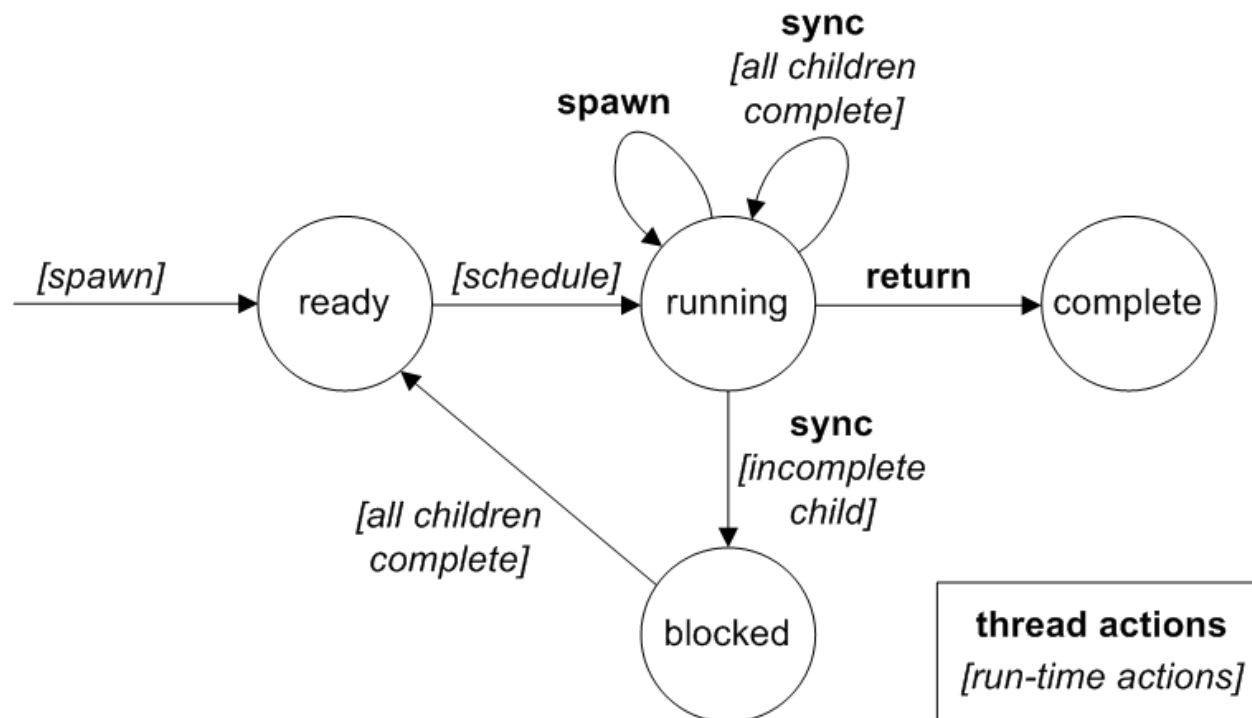
```
my_class x;
MyRefFunc(x); // pass by ref
```

# Const methods

- Methods with the `const` modifier cannot change the object
  - Cannot modify the internal state of the object
  - Object is “the same” after any const method is called
- References with the `const` modifier cannot be changed
  - Can only call const methods and read object properties
  - Only methods marked const are considered const methods
    - Methods which don’t change state are “non-const” unless marked
- Const references are very useful
  - Pass by reference makes it cheap to pass object to function
  - Const-ness means object can be safely used in parallel

# Recall: *Life-cycle of a task*

- Life-cycle of task due to interaction between task and run-time
  - Individual task calls **spawn**, **sync**, **return**
  - Cilk run-time must keep track of a task's children (*dependencies*)



# Scheduling through reference counts

- Each task has a reference count and a successor task
- The reference count identifies whether a task is blocked
  - If the reference count is zero then the task could be run
  - But only if it has been given to the task scheduler
  - Legal to create a task and not give it to the scheduler
  - *Note the difference: “reference count” vs “C++ reference”*

# Scheduling through reference counts

- Each task has a reference count and a successor task
- The reference count identifies whether a task is blocked
  - If the reference count is zero then the task could be run
  - But only if it has been given to the task scheduler
  - Legal to create a task and not give it to the scheduler
  - *Note the difference: “reference count” vs “C++ reference”*
- The successor task identifies the task blocked by this task
  - Same concept as “parent” in Cilk, but slightly more general
  - When a task completes it decrements the count of its successor

```
cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}
```

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

```

```

cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}

```

```

tbb::task * execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    spawn(t1);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
}
};

```

```

void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start, end);
    tbb::task::spawn_root_and_wait();
}

```

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

```

```

cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}

```

```

tbb::task * execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    spawn(t1);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
}
};

```

```

void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start, end);
    tbb::task::spawn_root_and_wait();
}

```



```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

```

```

    tbb::task * execute()
    {
        if(cond())
            return 0;
        set_ref_count(3);
        MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
        spawn(t1);
        MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
        spawn(t2);
        DoSomethingFirst();
        wait_for_all();
        DoSomethingElse();
    }
};

```

```

void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}

```

```

cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}

```

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

```

```

    tbb::task * execute()
    {
        if(cond())
            return 0;
        set_ref_count(3);
        MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
        spawn(t1);
        MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
        spawn(t2);
        DoSomethingFirst();
        wait_for_all();
        DoSomethingElse();
    }
};

```

```

void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}

```

```

cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}

```

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

    tbb::task * execute()
    {
        if(cond())
            return 0;
        set_ref_count(3);
        MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
        spawn(t1);
        MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
        spawn(t2);
        DoSomethingFirst();
        wait_for_all();
        DoSomethingElse();
    }
};

void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start, end);
    tbb::task::spawn_root_and_wait();
}

```

```

cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}

```

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

```

```

    tbb::task * execute()
    {
        if(cond())
            return 0;
        set_ref_count(3);
        MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
        spawn(t1);
        MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
        spawn(t2);
        DoSomethingFirst();
        wait_for_all();
        DoSomethingElse();
    }
};

```

```

void CreateTasks(int start, int end)

```

```

{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}

```

```

cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}

```

```

class MyTask
: public tbb::task
{
    int start, end;

    MyTask(int _start, int _end)
    { start=_start; end=_end; }

    tbb::task * execute()
    {
        if(cond())
            return 0;
        set_ref_count(3);
        MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
        spawn(t1);
        MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
        spawn(t2);
        DoSomethingFirst();
        wait_for_all();
        DoSomethingElse();
    }
};

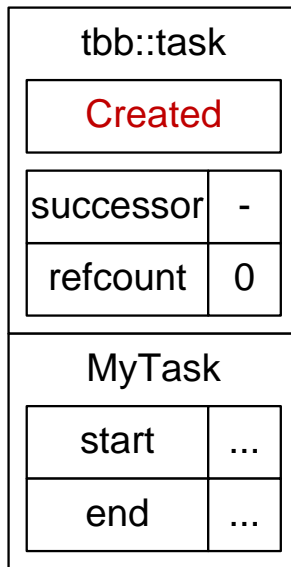
void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start, end);
    tbb::task::spawn_root_and_wait();
}

```

```

cilk void MyTask(int start, int end)
{
    if(cond())
        return 0;
    spawn MyTask(start, (start+end)/2);
    spawn MyTask((start+end)/2, end);
    DoSomethingFirst();
    sync;
    DoSomethingElse();
    return 0;
}

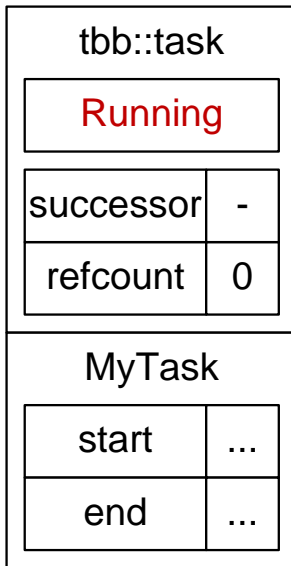
```




```

void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}

```



```
void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
```

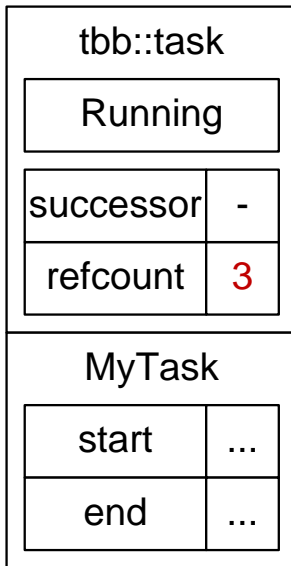


```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

```

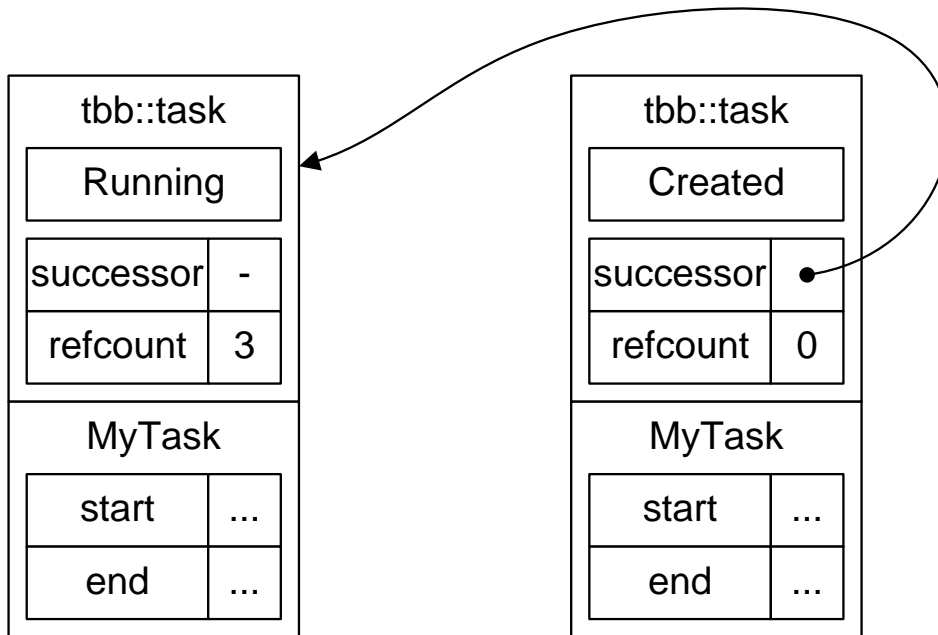





```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

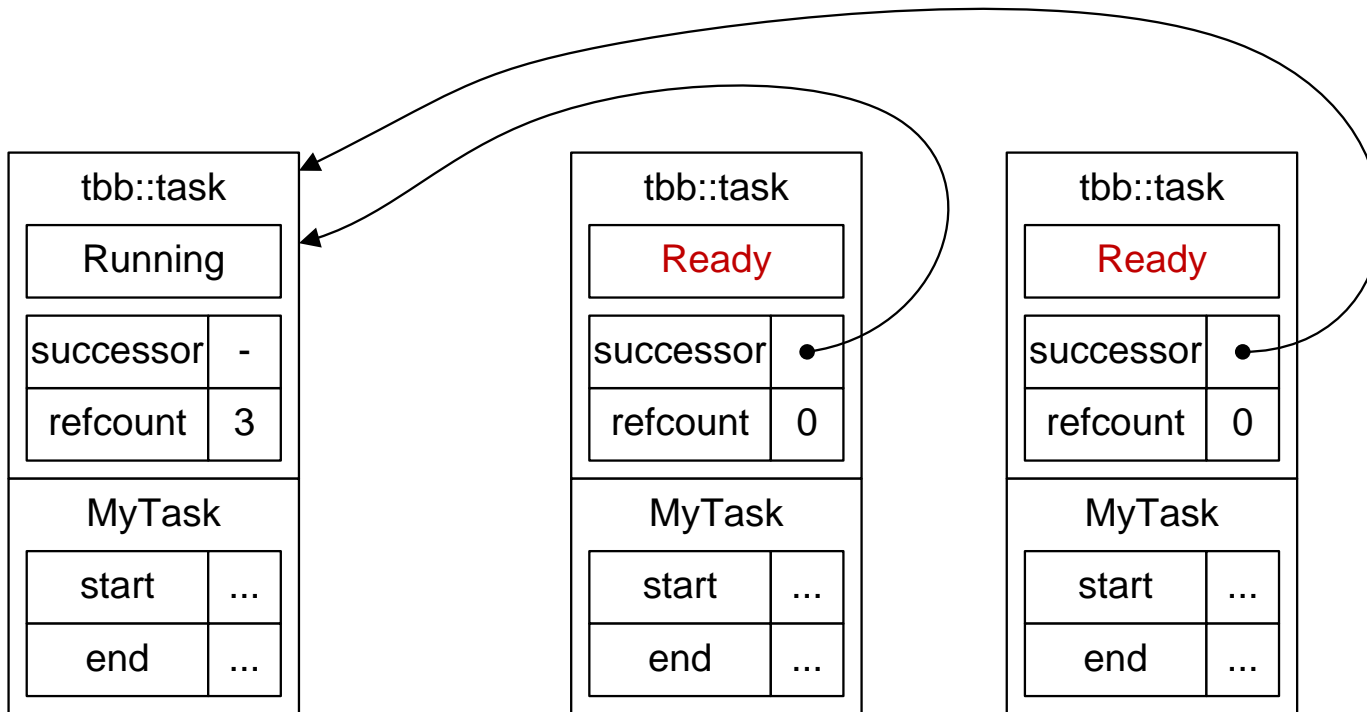
```



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```




```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

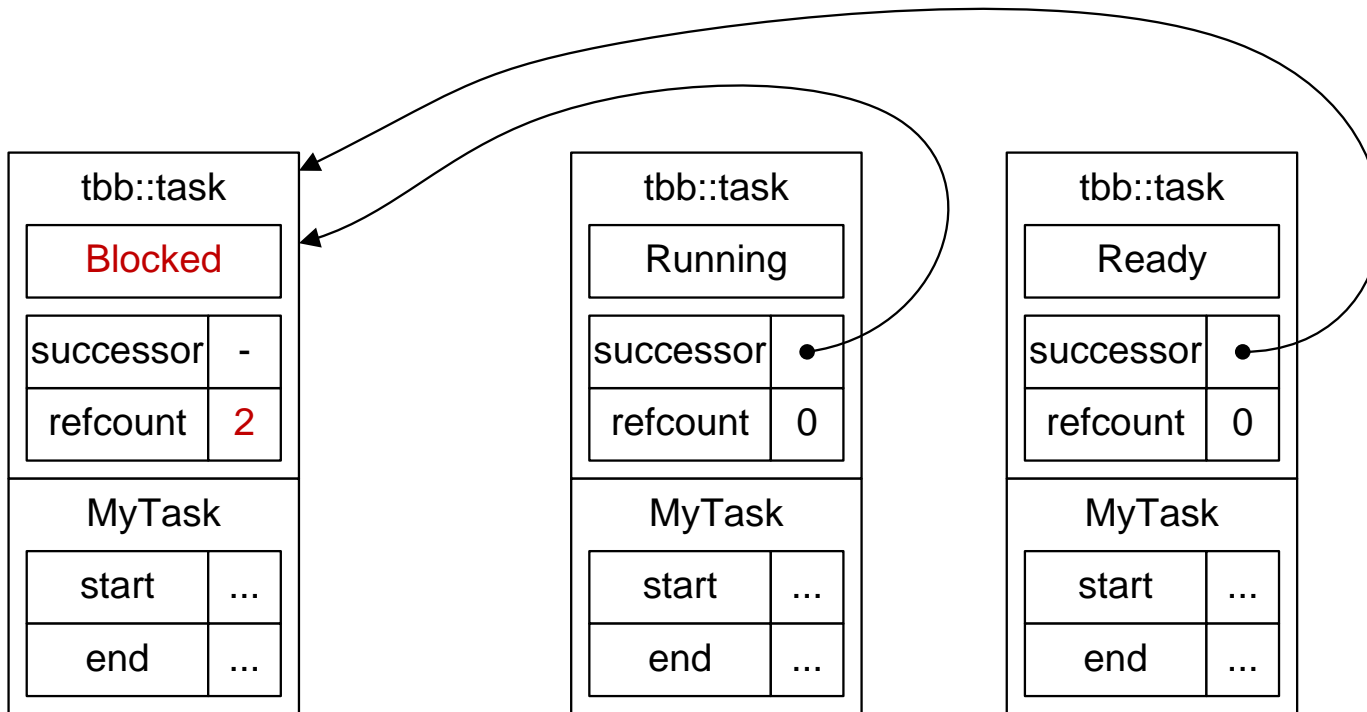
```



```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

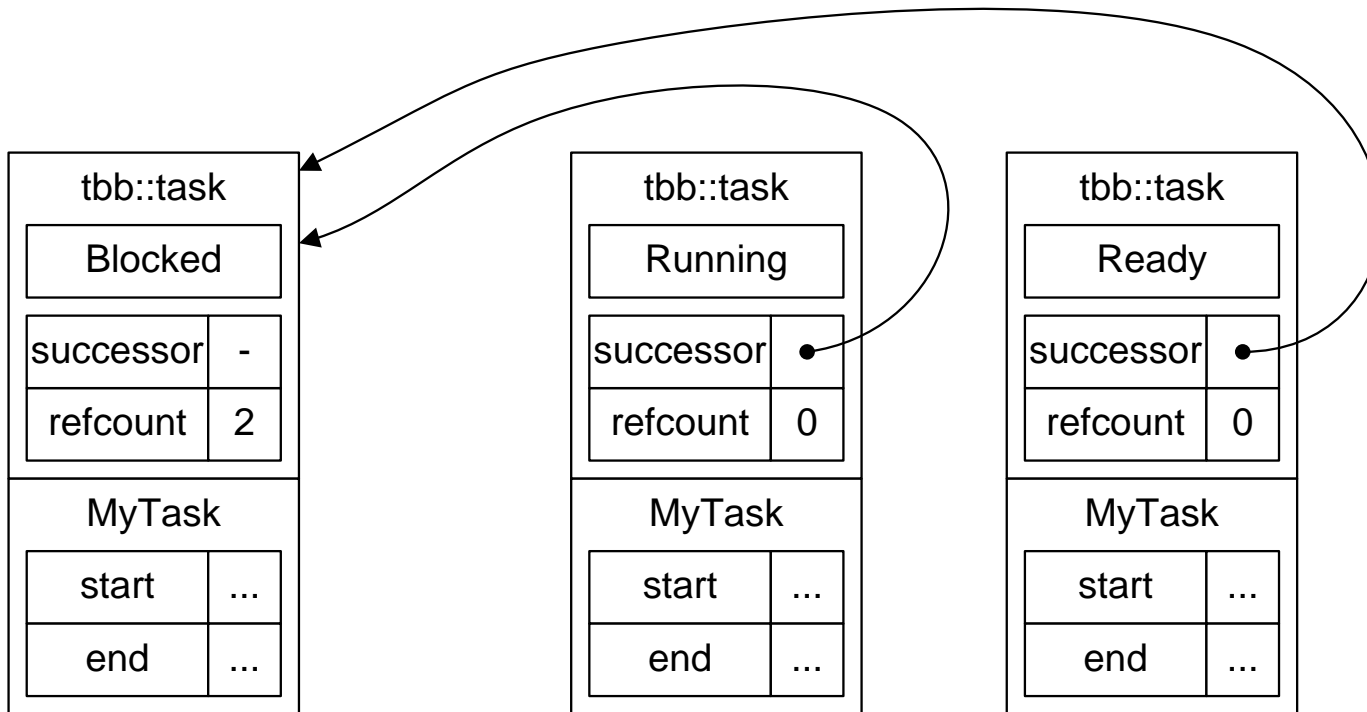
```



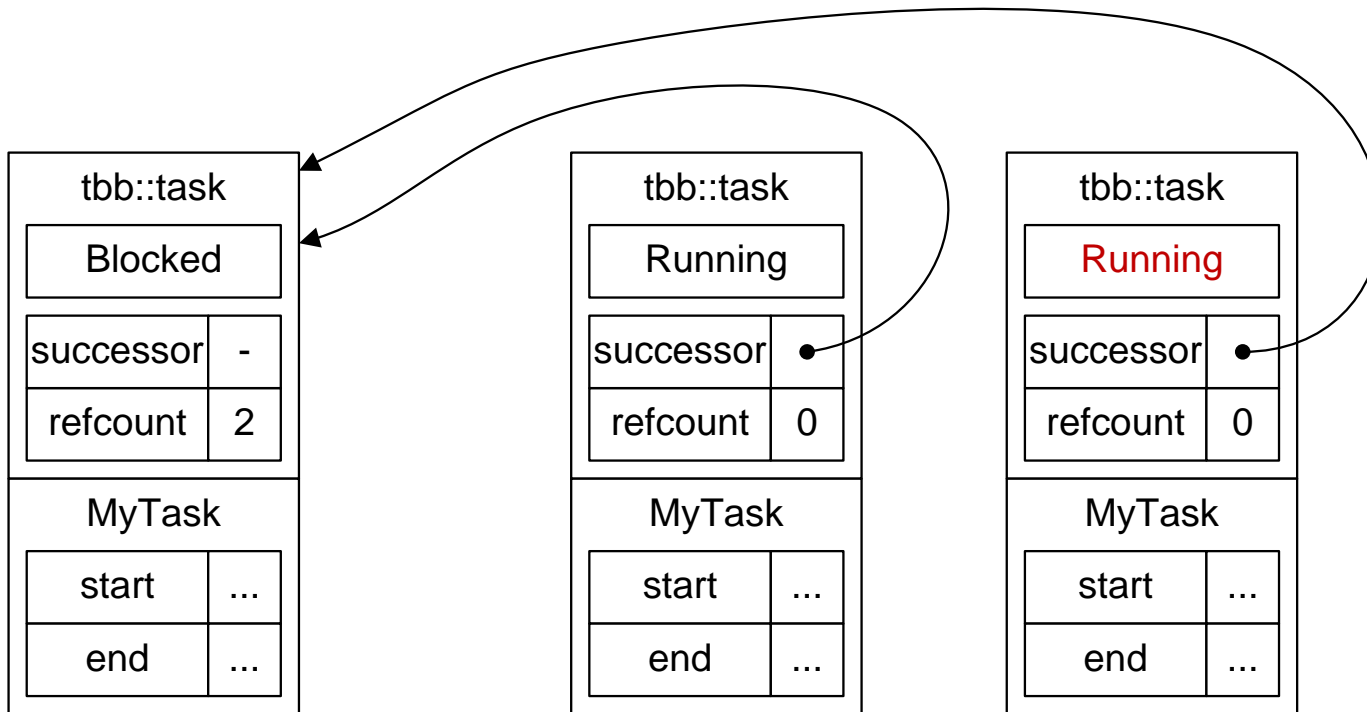
```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

```



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```

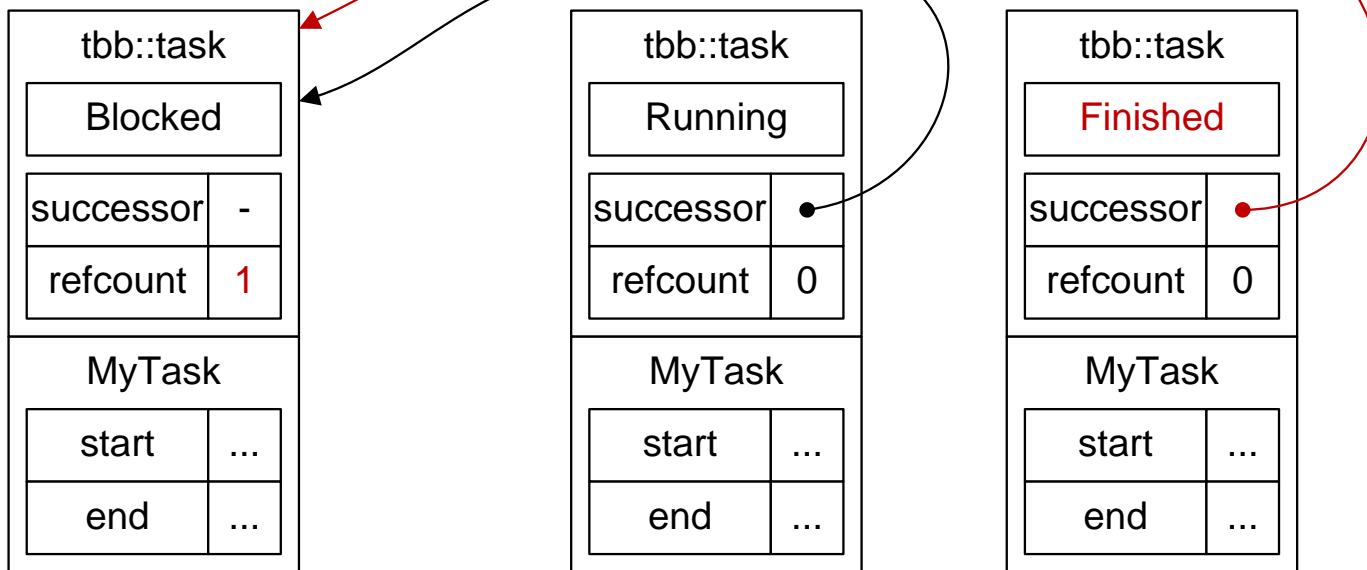


```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

```

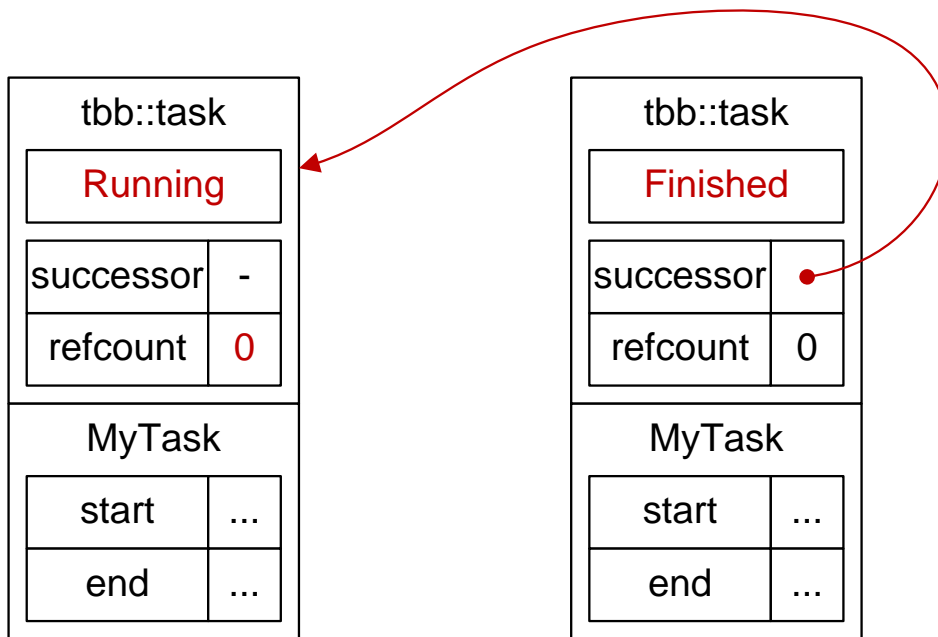





```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

```



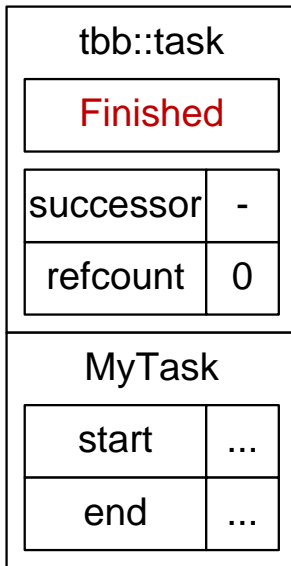
```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(3);
    MyTask &t1=*new(allocate_child()) MyTask(start,(start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}


```



```
void CreateTasks(int start, int end)
{
    MyTask &root=*new(allocate_root()) MyTask(start,end);
    tbb::task::spawn_root_and_wait();
}
```

# Managing reference counts

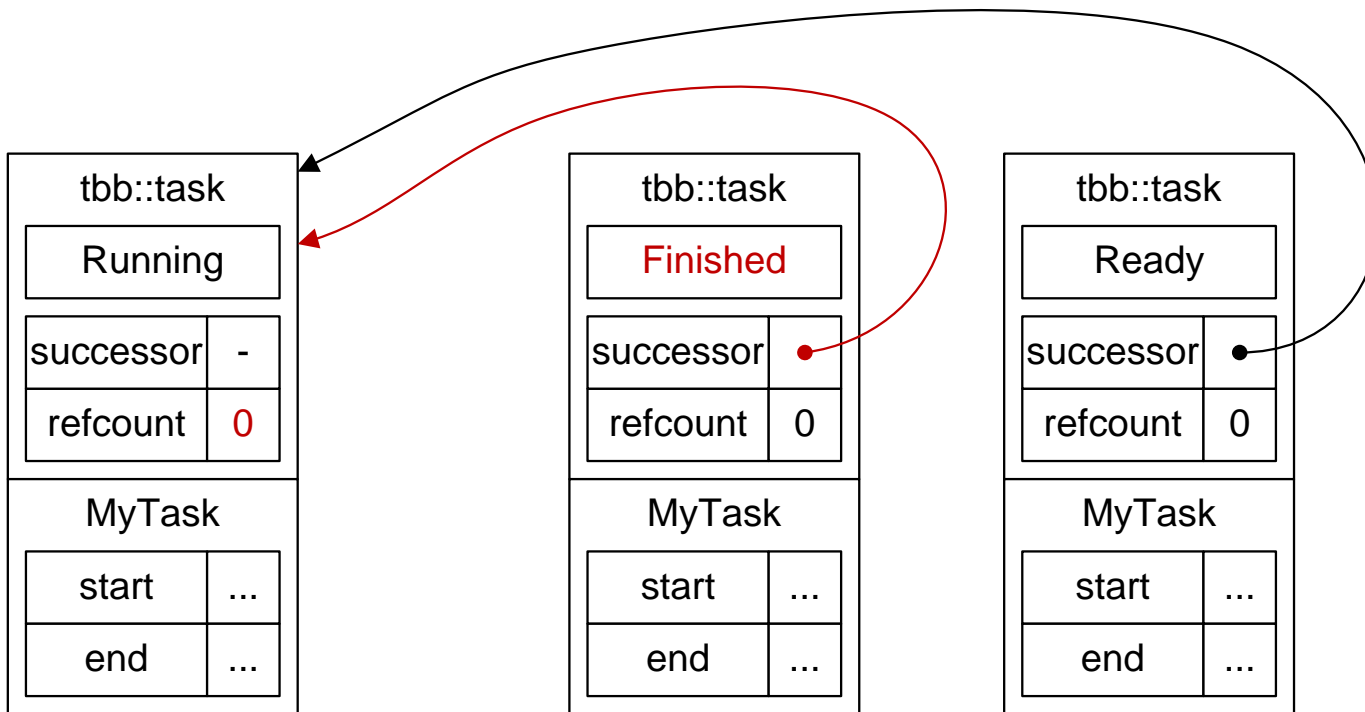
- What happens if we get the reference count wrong?
- Finishing task calls `decrement_ref_count` on successor
  - Automatically returns task to scheduler if count becomes zero



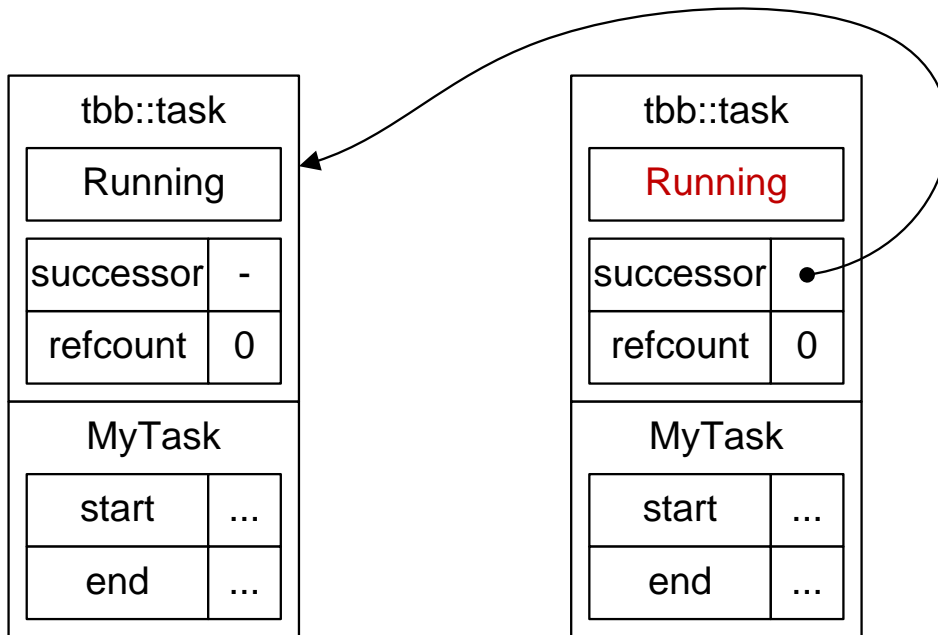
```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(1);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

```

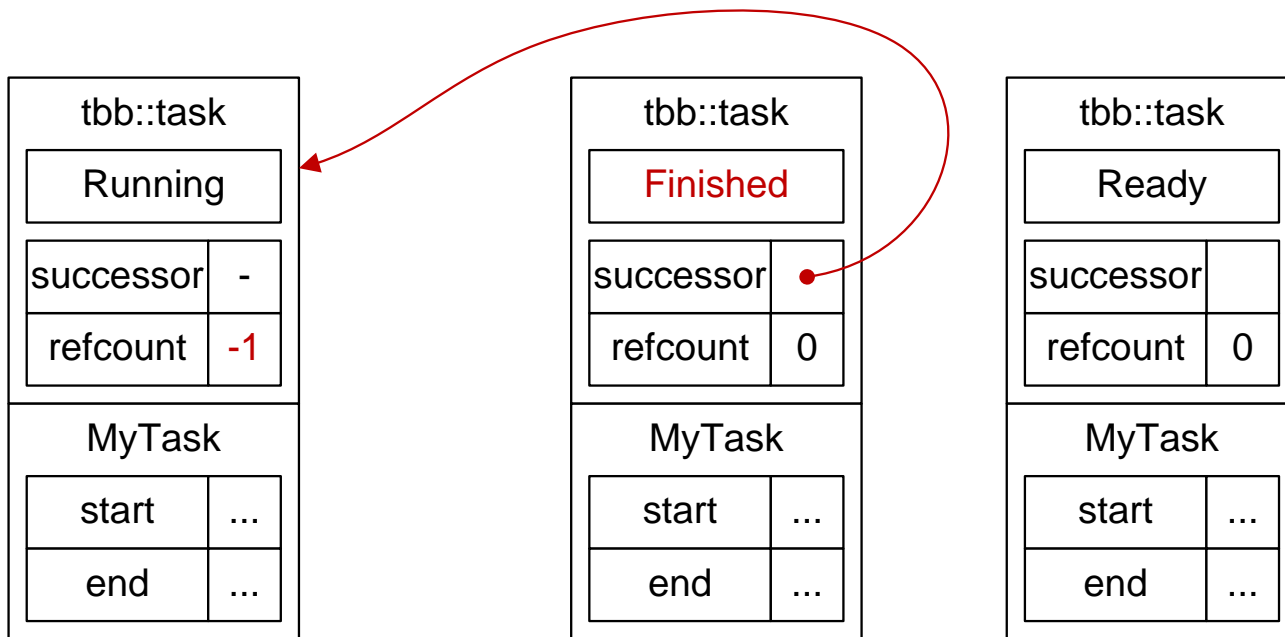


```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    set_ref_count(1);
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```



```
tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    set_ref_count(3);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}
```





```

tbb::task * MyTask::execute()
{
    if(cond())
        return 0;
    MyTask &t1=*new(allocate_child()) MyTask(start, (start+end)/2);
    MyTask &t2=*new(allocate_child()) MyTask((start+end)/2, end);
    spawn(t1);
    spawn(t2);
    set_ref_count(3);
    DoSomethingFirst();
    wait_for_all();
    DoSomethingElse();
    return 0;
}

```

# Some help is available

- TBB library comes in two forms: debug and release
  - release library does no error checking – all about speed
  - debug library will check reference counts at many points
- Choose library version at compilation and link stages
  - Debug: `#define TBB_USE_DEBUG=1` when compiling
    - On microsoft compilers it will automatically link the correct library
    - On other compilers use “`-ltbb`” vs “`-ltbb_debug`”
  - Usually maintain different release and debug settings
    - Debug: `/DTBB_USE_DEBUG=1 /MDd`
    - Release: `/DNDEBUG=1 /O2`
  - Can setup in Visual Studio or in a makefile

# Data-parallelism vs task parallelism

- Two very broad types of parallelism we've seen so far
- **Data-parallelism:** *do the same task lots of time in any order*
  - The code for the task stays the same for each execution
  - The input to the task changes with each execution
  - There are no dependencies between different executions
    - Often applied to elements of an array
    - Also described as “loop-parallelism”
- **Task-parallelism:** *do many different tasks with dependencies*
  - Each task has zero or more dependencies that must be met
  - Different tasks **may** have different code
  - More powerful than data-parallelism?